

# MA22M006\_MLP\_1

September 3, 2023

Here I will create a Neural Network and train using “back propogation” to recognise the hand written digits using MNIST Dataset. Taken Some Ideas rather Code for last leg of this assignment from [Here](#). Also Uses ChatGPT for plotting the confusion matrix, Surprisingly it has given the accurate response. Also I would like to thank **Rohan (MA22M025)** for helping me in backpropagation part.

## 1 Importing The Usefull Libraries

```
[ ]: #Importing the libraries that will be helpful during calculation
import numpy as np
import pandas as pd
import torch
import matplotlib.pyplot as plt
from torchvision.transforms import ToTensor
```

## 2 Loading DataSet From Pytorch

```
[ ]: from torchvision import datasets
training_data = datasets.MNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor() # it converts PIL image into pytorch tensor and also
    ↪normalise it
)

test_data = datasets.MNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor() # it converts PIL image into pytorch tensor and also
    ↪normalise it
)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>  
Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to

```

data/MNIST/raw/train-images-idx3-ubyte.gz
100%|      | 9912422/9912422 [00:00<00:00, 108961875.77it/s]
Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|      | 28881/28881 [00:00<00:00, 23237232.65it/s]
Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|      | 1648877/1648877 [00:00<00:00, 30874101.67it/s]
Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|      | 4542/4542 [00:00<00:00, 7153784.74it/s]
Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

```

## 2.1 Preprocessing The DATA

```

[ ]: VALID_RATIO = 0.9
    # Spilliting the data into train and validation sets
    n_train_examples = int(len(training_data) * VALID_RATIO)
    n_valid_examples = len(training_data) - n_train_examples
    train_data, valid_data = torch.utils.data.random_split(training_data,
                                                            [n_train_examples, n_valid_examples])

[ ]: def Data(Train,Valid,Test,Batch_Size=64):
    train_iterator = torch.utils.data.DataLoader(Train,
                                                shuffle=True,
                                                batch_size=Batch_Size,
                                                drop_last=True)

    valid_iterator =torch.utils.data.DataLoader(Valid,
                                                batch_size=6000)

    test_iterator =torch.utils.data.DataLoader(Test,

```

```

        batch_size=10000)
    return train_iterator,valid_iterator,test_iterator

```

```

[ ]: train_iterator,valid_iterator,test_iterator = Data(
    ↪Data(train_data,valid_data,test_data)

```

### 2.1.1 Visualisation of data

```

[ ]: for data in train_iterator:
    print(data[0].shape,data[1].shape)
    break # It contains 60000 images

```

torch.Size([64, 1, 28, 28]) torch.Size([64])

```

[ ]: for data in valid_iterator:
    print(data[0].shape,data[1])
    break # It contains 60000 images


```

torch.Size([6000, 1, 28, 28]) tensor([3, 6, 0, ..., 2, 9, 6])

```


[ ]: fig,ax= plt.subplots(nrows=4,ncols=4,figsize=(16,8),sharex=True,sharey=True)
    for i in range(4):
        for j in range(4):
            k=(i+1)*(j+1)+100
            ax[i,j].imshow(training_data.data[k],cmap='gray')
            ax[i,j].set_title(f'True Label: {training_data.targets[k]}')
    plt.tight_layout()

```



### 2.1.2 Distribution of Training Data into different classes. We can see data is uniformly distributed

```
[ ]: plt.hist(train_iterator.dataset.dataset.targets,bins=40)
plt.show()
```



## 3 Changing the dimensions of data in order to use it for training.

```
[ ]: # Creates one hot vectors out of given targets
def one_hot(Y):
    one_hot_Y = torch.zeros(len(Y), 10)
    one_hot_Y[range(len(Y)), Y] = 1
    one_hot_Y = one_hot_Y.T
    return one_hot_Y    # 10X64
```

```
[ ]: # Flattening the images for training the model
X_train=[]
Y_train=[]
```

```

for X, Y in train_iterator:
    X=torch.flatten(X,start_dim=1)
    Y=one_hot(Y) # 10X 64
    X_train.append(X)
    Y_train.append(Y)

```

```
[ ]: X_train[0].shape
```

```
[ ]: torch.Size([64, 784])
```

```
[ ]: Y_train[0].shape
```

```
[ ]: torch.Size([10, 64])
```

```
[ ]: X_valid = [torch.flatten(x[0],start_dim=1) for x in valid_iterator]
      Y_valid = [y[1] for y in valid_iterator]
```

```
[ ]: X_valid[0].shape
```

```
[ ]: torch.Size([6000, 784])
```

```
[ ]: X_test = [torch.flatten(x[0],start_dim=1) for x in test_iterator]
      Y_test = [y[1] for y in test_iterator]
```

```
[ ]: print(X_test[0].shape,Y_test[0].shape)
```

```
torch.Size([10000, 784]) torch.Size([10000])
```

### 3.1 Loss Function:

```

[ ]: def Cross_entropy(Y, pred,lambda1=0,params_sum=0):
      J = Y * pred
      J=torch.sum(J,0)
      return -torch.sum(torch.log(J))+ lambda1*params_sum

```

#### 3.1.1 Activation Functions

Relu Activation Function:

$$ReLU(x) = \max(0, x)$$

```

[ ]: def ReLU(x):
      return torch.maximum(torch.zeros_like(x),x)

def ReLU_deriv(Z):
    return Z > 0

```

Sigmoid Activation Function:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

```
[ ]: def Sigmoid(x):  
    return 1/(1+torch.exp(-x))
```

```
[ ]: def Grad_Sigmoid(x):  
    return Sigmoid(x)*(1-Sigmoid(x))
```

Tanh Activation Function:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

```
[ ]: def Tanh(x):  
    return torch.div((torch.exp(x)-torch.exp(-x)),(torch.exp(x)+torch.exp(-x)))
```

```
[ ]: def Grad_Tanh(x):  
    return 1-Tanh(x)*Tanh(x)
```

Softmax Function:

$$\text{Softmax}(X_{n \times 1}) = \left( \frac{e^{x_i}}{\sum_j e^j} \right)_{n \times 1}$$

```
[ ]: def Softmax(X,dim=0):  
    Y=torch.zeros_like(X)  
    if dim==0:  
        i=0  
        for x in X.T:  
            Y.T[i]=torch.exp(x)/sum(torch.exp(x))  
            i+=1  
    else:  
        i=0  
        for x in X:  
            Y[i]=torch.exp(x)/sum(torch.exp(x))  
            i+=1  
    return Y
```

## 4 The Model

### 4.1 Defining The Architecture

```
[ ]: torch.seed()  
def init_params(In = [784,500,250,100],Out = [500,250,100,10]):  
    param=[]  
    for I,0 in zip(In,Out):  
        M=np.sqrt(6/(I+0))
```

```

W=torch.FloatTensor(0, I).uniform_(-M,M)
b=torch.FloatTensor(0,1).uniform_(-M,M)
param.append(W)
param.append(b)
return param

```

```

[ ]: params=init_params()
for a in params:
    print(a.shape)

```

```

torch.Size([500, 784])
torch.Size([500, 1])
torch.Size([250, 500])
torch.Size([250, 1])
torch.Size([100, 250])
torch.Size([100, 1])
torch.Size([10, 100])
torch.Size([10, 1])

```

#### 4.1.1 Forward Propagation

```

[ ]: def forward_prop(param, X,Activation=Sigmoid,OutAct=Softmax):
    ↪#hiddenacti,outacti
    #print('X',X)
    Z1=torch.mm(param[0],X.T)+param[1]
    #print('Z1',Z1)
    Activations=[X.T, Z1,Activation(Z1)]
    #print('A1',Activations[-1])
    for i in range(2,len(param),2):
        Z = torch.mm(param[i],Activations[-1])+param[i+1]
        Activations.append(Z)
        #print('Z',Z)
        if i!=len(param)-2:
            Activations.append(Activation(Z))
            # print("A",torch.nn.functional.sigmoid(Z))
        else:
            Activations.append(OutAct(Z,dim=0))
            #print("A",torch.nn.functional.softmax(Z,dim=0))
    return Activations

```

```

[ ]: Activ=forward_prop(params,X_train[0])

```

```

[ ]: for a in Activ:
    print(a.shape)

```

```

torch.Size([784, 64])
torch.Size([500, 64])

```

```

torch.Size([500, 64])
torch.Size([250, 64])
torch.Size([250, 64])
torch.Size([100, 64])
torch.Size([100, 64])
torch.Size([10, 64])
torch.Size([10, 64])

```

```
[ ]: Activ[-1].T[0]  # Looking into output from softmax
```

```
[ ]: tensor([0.0832, 0.0293, 0.1133, 0.1879, 0.1275, 0.1134, 0.0691, 0.0290, 0.1510,
           0.0964])
```

#### 4.1.2 Back Ward Propagation

```
[ ]: def backward_pass(X, Y, activations, parameters, grad_act = Grad_Sigmoid):
    grads = {}
    m = X.shape[1]
    dZ = activations[-1] - Y  #10X64
    dZ=dZ/m
    #print(dZ.shape)
    #print(activations[-3].shape)
    grads[f'db{len(parameters) // 2}'] = torch.sum(dZ, dim=1, keepdim=True)
    grads[f'dW{len(parameters) // 2}'] = torch.mm(dZ, activations[-3].T)

    for i in reversed(range(1, len(parameters) // 2)):
        dZ = torch.mm(parameters[2*i].T, dZ) *(grad_act(activations[2*i]))

        grads[f'db{i}'] = torch.sum(dZ, dim=1, keepdim=True)
        grads[f'dW{i}'] = torch.mm(dZ, activations[2*i-2].T)
    return grads
```

```
[ ]: grad=backward_pass(X_train[0],Y_train[0],Activ,params)
```

```
[ ]: for a in reversed(grad):
    print(a)
    print(grad[a].shape)
```

```

dW1
torch.Size([500, 784])
db1
torch.Size([500, 1])
dW2
torch.Size([250, 500])
db2
torch.Size([250, 1])
dW3

```



```

torch.Size([100, 250])
db3
torch.Size([100, 1])
dW4
torch.Size([10, 100])
db4
torch.Size([10, 1])

```

### 4.1.3 Updation of Parameters

```

[ ]: def update_params(param,gradp, alpha, lambda1=0):
    r=0
    for i ,j in zip(param,reversed(gradp)):
        #print(i.shape,gradp[j].shape)
        if r%2+1:
            param[r]=(1-lambda1*alpha)*i-alpha*gradp[j]
        else:
            param[r]=i-alpha*gradp[j]
        r+=1
    return param

```

```

[ ]: uparam=update_params(params,grad,0.01)

```

### 4.1.4 Predictions and Accuracy functions

```

[ ]: def get_predictions(A4):
    return torch.argmax(A4, 0)

def get_accuracy(predictions, Y):
    #print(predictions, Y)
    X=torch.argmax(Y,0)
    #print(X.shape)
    #print(predictions.shape)
    return torch.sum(predictions == X) / len(X)

```

## 4.2 Gradient Descent Algorithm

```

[ ]: def gradient_descent(X_train,Y_train, X_test, Y_test , X_valid, Y_valid,
    ↪Activation=Sigmoid,
    grad_act=Grad_Sigmoid, OutAct = Softmax, alpha=0.01,
    ↪epoch= 20,lambda1=0):
    params = init_params()
    #print('weight_Initialised.....')
    EpochLoss_train=[]
    EpochLoss_valid=[]
    for epo in range(epoch):

```

```

for X,Y in zip(X_train,Y_train):

    activations = forward_prop(params,X, Activation, OutAct)

    gradp = backward_pass(X,Y, activations, params,grad_act)

    params = update_params(params, gradp, alpha,lambda1)

    #predictions = get_predictions(activations[-1])
    param_sum = sum([torch.sum(params[i]) for i in range(0,8,2)])
    #EpochLoss_train.
    ↪append(Cross_entropy(Y,activations[-1],lambda1,params_sum=param_sum))

    Activ1=forward_prop(params, X_valid[0], Activation, OutAct)

    EpochLoss_valid.append(Cross_entropy(one_hot(Y_valid[0]),Activ1[-1])/6000)
    train_true=0
    loss=0
    for X,Y in zip(X_train,Y_train):

        activt=forward_prop(params, X, Activation, OutAct)
        Y_hat_train = get_predictions(activt[-1])
        loss += Cross_entropy(Y,activt[-1],lambda1,params_sum=param_sum)
        train_true += torch.sum(Y_hat_train==torch.argmax(Y,0))

    EpochLoss_train.append(loss/54000)
    #activations = forward_prop(params, X)
    Y_hat_val=get_predictions(Activ1[-1])
    #Y_hat_train = get_predictions(activations[-1])
    print(f'***** Epoch {epo+1}␣
    ↪*****')
    print('Training Accuracy: ',(train_true/54000).item()*100)
    print('Validation Accuracy: ',(torch.sum(Y_hat_val==Y_valid[0])/6000).
    ↪item()*100)

    ␣
    ↪print('-----')
    Activ2=forward_prop(params, X_test[0], Activation, OutAct)
    Y_hat_test = get_predictions(Activ2[-1])
    print('Test Accuracy: ',(torch.sum(Y_hat_test==Y_test[0])/10000).item()*100)
    return params,EpochLoss_train,EpochLoss_valid , Y_hat_test

```

### 4.3 Confusion Matrix

```
[ ]: def Confusion_Matrix(Y_pre, Y):  
    number_of_classes = 10  
    conf_matrix = np.zeros((number_of_classes, number_of_classes), dtype=int)  
    for true_label, prediction in zip(Y_pre, Y):  
        conf_matrix[true_label, prediction] += 1  
    return conf_matrix
```

```
[ ]: def plot_confusion_matrix(confusion_matrix):  
    plt.imshow(confusion_matrix, interpolation='nearest', cmap=plt.cm.Blues)  
    plt.title("Confusion Matrix")  
    plt.colorbar()  
  
    num_classes = confusion_matrix.shape[0]  
  
    plt.xticks(np.arange(num_classes), range(num_classes))  
    plt.yticks(np.arange(num_classes), range(num_classes))  
  
    plt.ylabel('True label')  
    plt.xlabel('Predicted label')  
  
    # Add text annotations for each cell  
    for i in range(num_classes):  
        for j in range(num_classes):  
            plt.text(j, i, str(confusion_matrix[i, j]),  
                    horizontalalignment="center", color="black")  
  
    plt.show()
```

### 4.4 (1) Using Sigmoid and softmax

Here I am using sigmoid activation in hidden layers and softmax at the output layer doing gradient descent with learning rate 1 and runs for 15 epoch

```
[ ]: # For differen learning rate please change alpha  
Params = gradient_descent(X_train, Y_train, X_test, Y_test, X_valid, Y_valid,  
    ↪ Activation=Sigmoid,  
    grad_act=Grad_Sigmoid, OutAct = Softmax, alpha=1, epoch=  
    ↪ 15)
```

\*\*\*\*\* Epoch 1 \*\*\*\*\*

Training Accuracy: 41.93703830242157

Validation Accuracy: 41.11666679382324

\*\*\*\*\* Epoch 2 \*\*\*\*\*

Training Accuracy: 78.2537043094635

Validation Accuracy: 78.29999923706055

```

-----
***** Epoch 3 *****
Training Accuracy: 84.62036848068237
Validation Accuracy: 84.26666855812073
-----
***** Epoch 4 *****
Training Accuracy: 87.37407326698303
Validation Accuracy: 87.23333477973938
-----
***** Epoch 5 *****
Training Accuracy: 88.80741000175476
Validation Accuracy: 88.40000033378601
-----
***** Epoch 6 *****
Training Accuracy: 89.7166693687439
Validation Accuracy: 89.13333415985107
-----
***** Epoch 7 *****
Training Accuracy: 90.39815068244934
Validation Accuracy: 89.81666564941406
-----
***** Epoch 8 *****
Training Accuracy: 90.82221984863281
Validation Accuracy: 90.16666412353516
-----
***** Epoch 9 *****
Training Accuracy: 91.15926027297974
Validation Accuracy: 90.56666493415833
-----
***** Epoch 10 *****
Training Accuracy: 91.39629602432251
Validation Accuracy: 90.88333249092102
-----
***** Epoch 11 *****
Training Accuracy: 91.61296486854553
Validation Accuracy: 90.9666657447815
-----
***** Epoch 12 *****
Training Accuracy: 91.79444313049316
Validation Accuracy: 91.18333458900452
-----
***** Epoch 13 *****
Training Accuracy: 91.95926189422607
Validation Accuracy: 91.33333563804626
-----
***** Epoch 14 *****
Training Accuracy: 92.08333492279053
Validation Accuracy: 91.56666398048401

```

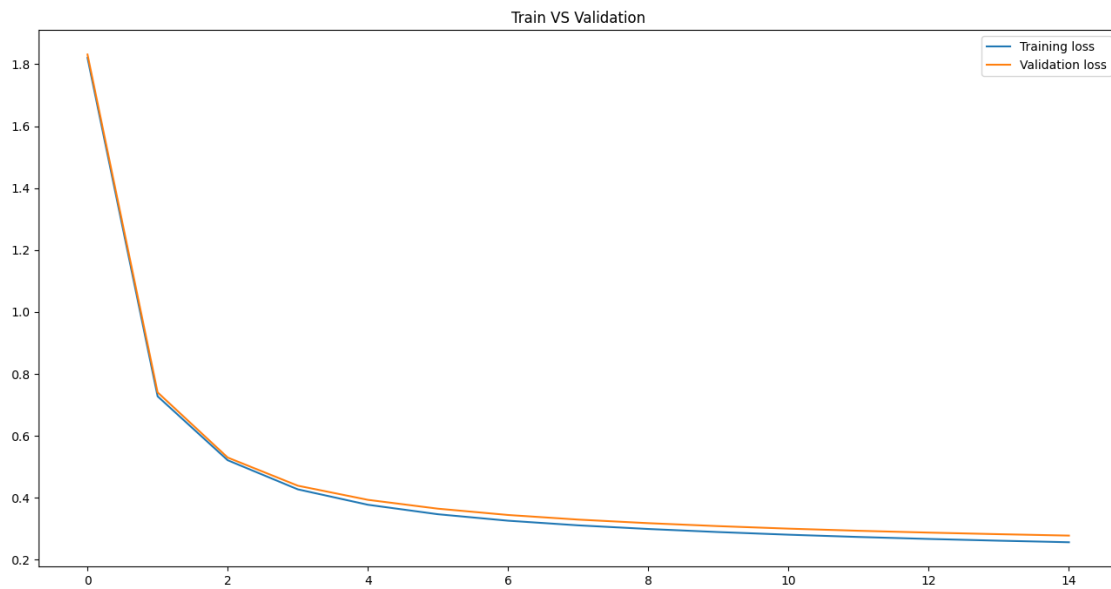
-----  
\*\*\*\*\* Epoch 15 \*\*\*\*\*

Training Accuracy: 92.25555658340454


Validation Accuracy: 91.60000085830688  
-----

Test Accuracy: 92.3799991607666

```
[ ]: plt.figure(figsize=(16,8))  
plt.plot(Params[1],label='Training loss')  
plt.plot(Params[2],label='Validation loss')  
plt.title('Train VS Validation')  
plt.legend()  
plt.show()
```



```
[ ]: Mat = Confusion_Matrix(Params[-1],Y_test[0])  
plot_confusion_matrix(Mat)
```



#### 4.5 (1) Using ReLU and softmax

Here I am using ReLU activation in hidden layers and softmax at the output layer doing gradient descent with learning rate 0.1 and runs for 15 epoch

```
[ ]: Params = gradient_descent(X_train,Y_train, X_test, Y_test , X_valid, Y_valid,
    ↪Activation=ReLU,
                                grad_act=ReLU_deriv, OutAct = Softmax,  alpha=1, epoch= 15)
```

```
***** Epoch 1 *****
Training Accuracy:  95.2129602432251
Validation Accuracy:  94.81666684150696
-----
```

```
***** Epoch 2 *****
Training Accuracy:  97.25925922393799
Validation Accuracy:  96.66666388511658
-----
```

```
***** Epoch 3 *****
Training Accuracy:  98.14444184303284
Validation Accuracy:  97.2166657447815
```

```


-----
***** Epoch 4 *****
Training Accuracy: 98.56851696968079
Validation Accuracy: 97.46666550636292
-----
***** Epoch 5 *****
Training Accuracy: 98.78518581390381
Validation Accuracy: 97.38333225250244
-----
***** Epoch 6 *****
Training Accuracy: 99.14259314537048
Validation Accuracy: 97.53333330154419
-----
***** Epoch 7 *****
Training Accuracy: 99.33518767356873
Validation Accuracy: 97.60000109672546
-----
***** Epoch 8 *****
Training Accuracy: 99.4962990283966
Validation Accuracy: 97.61666655540466
-----
***** Epoch 9 *****
Training Accuracy: 99.58703517913818
Validation Accuracy: 97.66666889190674
-----
***** Epoch 10 *****
Training Accuracy: 99.65925812721252
Validation Accuracy: 97.56666421890259
-----
***** Epoch 11 *****
Training Accuracy: 99.71296191215515
Validation Accuracy: 97.66666889190674
-----
***** Epoch 12 *****
Training Accuracy: 99.77592825889587
Validation Accuracy: 97.63333201408386
-----
***** Epoch 13 *****
Training Accuracy: 99.83333349227905
Validation Accuracy: 97.73333072662354
-----
***** Epoch 14 *****
Training Accuracy: 99.86110925674438
Validation Accuracy: 97.73333072662354
-----
***** Epoch 15 *****
Training Accuracy: 99.88333582878113
Validation Accuracy: 97.75000214576721

```

---


Test Accuracy: 98.07000160217285

```
[ ]: plt.figure(figsize=(16,8))
plt.plot(Params[1],label='Training loss')
plt.plot(Params[2],label='Validation loss')
plt.title('Train VS Validation')
plt.legend()
plt.show()
```



```
[ ]: Mat = Confusion_Matrix(Params[-1],Y_test[0])
plot_confusion_matrix(Mat)
```





[ ]:

#### 4.6 (1) Using Tanh and softmax

Here I am using Tanh activation in hidden layers and softmax at the output layer doing gradient descent with learning rate 0.1 and runs for 15 epoch

```
[ ]: Params = gradient_descent(X_train,Y_train, X_test, Y_test , X_valid, Y_valid,
    ↪Activation=Tanh,
                                grad_act=Grad_Tanh, OutAct = Softmax, alpha=0.1, epoch=
    ↪15)
```

\*\*\*\*\* Epoch 1 \*\*\*\*\*

Training Accuracy: 89.08148407936096

Validation Accuracy: 88.84999752044678

\*\*\*\*\* Epoch 2 \*\*\*\*\*

Training Accuracy: 90.65740704536438

Validation Accuracy: 90.6499981880188

```

***** Epoch 3 *****
Training Accuracy: 91.55185222625732
Validation Accuracy: 91.50000214576721
-----
***** Epoch 4 *****
Training Accuracy: 92.08518266677856
Validation Accuracy: 91.81666374206543
-----
***** Epoch 5 *****
Training Accuracy: 92.57592558860779
Validation Accuracy: 92.1999990940094
-----
***** Epoch 6 *****
Training Accuracy: 92.91852116584778
Validation Accuracy: 92.58333444595337
-----
***** Epoch 7 *****
Training Accuracy: 93.20926070213318
Validation Accuracy: 92.91666746139526
-----
***** Epoch 8 *****
Training Accuracy: 93.47962737083435
Validation Accuracy: 93.08333396911621
-----
***** Epoch 9 *****
Training Accuracy: 93.73703598976135
Validation Accuracy: 93.31666827201843
-----
***** Epoch 10 *****
Training Accuracy: 93.99259090423584
Validation Accuracy: 93.56666803359985
-----
***** Epoch 11 *****
Training Accuracy: 94.24999952316284
Validation Accuracy: 93.69999766349792
-----
***** Epoch 12 *****
Training Accuracy: 94.45740580558777
Validation Accuracy: 93.94999742507935
-----
***** Epoch 13 *****
Training Accuracy: 94.67407464981079
Validation Accuracy: 94.0999984741211
-----
***** Epoch 14 *****
Training Accuracy: 94.862961769104
Validation Accuracy: 94.23333406448364
-----

```


\*\*\*\*\* Epoch 15 \*\*\*\*\*

Training Accuracy: 95.02037167549133


Validation Accuracy: 94.38333511352539

-----  
Test Accuracy: 94.4899762535095

```
[ ]: plt.figure(figsize=(16,8))
plt.plot(Params[1],label='Training loss')
plt.plot(Params[2],label='Validation loss')
plt.title('Train VS Validation')
plt.legend()
plt.show()
```



```
[ ]: Mat = Confusion_Matrix(Params[-1],Y_test[0])
plot_confusion_matrix(Mat)
```



## 5 With Regularisation

### 5.0.1 Sigmoid

```
[ ]: Params = gradient_descent(X_train,Y_train, X_test, Y_test , X_valid, Y_valid,
    ↳Activation=Sigmoid,
                                grad_act=Grad_Sigmoid, OutAct = Softmax, alpha=0.1,
    ↳epoch= 15,lambdai=0.01)
```

\*\*\*\*\* Epoch 1 \*\*\*\*\*

Training Accuracy: 11.299999803304672

Validation Accuracy: 10.649999976158142

\*\*\*\*\* Epoch 2 \*\*\*\*\*

Training Accuracy: 11.299999803304672

Validation Accuracy: 10.649999976158142

\*\*\*\*\* Epoch 3 \*\*\*\*\*

Training Accuracy: 11.299999803304672

Validation Accuracy: 10.649999976158142

```

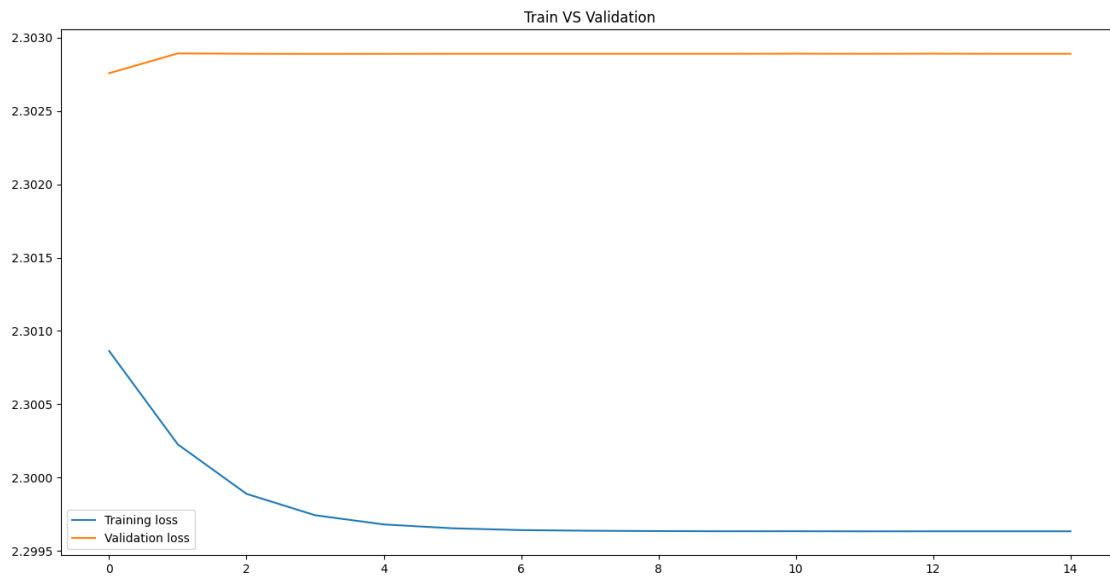
-----
***** Epoch 4 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 5 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 6 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 7 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 8 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 9 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 10 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 11 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 12 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 13 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 14 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142
-----
***** Epoch 15 *****
Training Accuracy: 11.299999803304672
Validation Accuracy: 10.649999976158142

```


---

Test Accuracy: 11.349999904632568

```
[ ]: plt.figure(figsize=(16,8))
plt.plot(Params[1],label='Training loss')
plt.plot(Params[2],label='Validation loss')
plt.title('Train VS Validation')
plt.legend()
plt.show()
```



```
[ ]: Mat = Confusion_Matrix(Params[-1],Y_test[0])
plot_confusion_matrix(Mat)
```



### 5.0.2 ReLU

```
[ ]: Params = gradient_descent(X_train,Y_train, X_test, Y_test , X_valid, Y_valid,
    ↪Activation=ReLU,
                                grad_act=ReLU_deriv, OutAct = Softmax, alpha=0.1, epoch=
    ↪15,lambda1=0.01)
```

```
***** Epoch 1 *****
Training Accuracy: 73.8388959884644
Validation Accuracy: 72.89999723434448
-----
***** Epoch 2 *****
Training Accuracy: 71.65370583534241
Validation Accuracy: 71.48333191871643
-----
***** Epoch 3 *****
Training Accuracy: 70.78889012336731
Validation Accuracy: 70.59999704360962
-----
***** Epoch 4 *****
```

Training Accuracy: 71.11851572990417  
Validation Accuracy: 70.86666822433472

-----  
\*\*\*\*\* Epoch 5 \*\*\*\*\*  
Training Accuracy: 71.59073948860168  
Validation Accuracy: 71.16666436195374  
-----

\*\*\*\*\* Epoch 6 \*\*\*\*\*  
Training Accuracy: 71.85184955596924  
Validation Accuracy: 71.5833306312561  
-----

\*\*\*\*\* Epoch 7 \*\*\*\*\*  
Training Accuracy: 72.03333377838135  
Validation Accuracy: 71.66666388511658  
-----

\*\*\*\*\* Epoch 8 \*\*\*\*\*  
Training Accuracy: 72.13703989982605  
Validation Accuracy: 71.63333296775818  
-----

\*\*\*\*\* Epoch 9 \*\*\*\*\*  
Training Accuracy: 72.2333312034607  
Validation Accuracy: 71.64999842643738  
-----

\*\*\*\*\* Epoch 10 \*\*\*\*\*  
Training Accuracy: 72.28147983551025  
Validation Accuracy: 71.73333168029785  
-----

\*\*\*\*\* Epoch 11 \*\*\*\*\*  
Training Accuracy: 72.303706407547  
Validation Accuracy: 71.73333168029785  
-----

\*\*\*\*\* Epoch 12 \*\*\*\*\*  
Training Accuracy: 72.32592701911926  
Validation Accuracy: 71.71666622161865  
-----

\*\*\*\*\* Epoch 13 \*\*\*\*\*  
Training Accuracy: 72.3370373249054  
Validation Accuracy: 71.70000076293945  
-----


\*\*\*\*\* Epoch 14 \*\*\*\*\*  
Training Accuracy: 72.32407331466675  
Validation Accuracy: 71.70000076293945  
-----

\*\*\*\*\* Epoch 15 \*\*\*\*\*  
Training Accuracy: 72.31296300888062  
Validation Accuracy: 71.70000076293945  
-----


Test Accuracy: 72.99000024795532



```
[ ]: plt.plot(figsize=(16,8))
plt.plot(Params[1],label='Training loss')
plt.plot(Params[2],label='Validation loss')
plt.title('Train VS Validation')
plt.legend()
plt.show()
```



```
[ ]: Mat = Confusion_Matrix(Params[-1],Y_test[0])
plot_confusion_matrix(Mat)
```



### 5.0.3 Tanh

```
[ ]: Params = gradient_descent(X_train,Y_train, X_test, Y_test , X_valid, Y_valid,
    ↪Activation=Tanh,
                                grad_act=Grad_Tanh, OutAct = Softmax, alpha=0.1, epoch=
    ↪15,lambda1=0.01)
```

\*\*\*\*\* Epoch 1 \*\*\*\*\*

Training Accuracy: 81.77777528762817

Validation Accuracy: 81.30000233650208

\*\*\*\*\* Epoch 2 \*\*\*\*\*

Training Accuracy: 75.61296224594116

Validation Accuracy: 75.3000020980835

\*\*\*\*\* Epoch 3 \*\*\*\*\*

Training Accuracy: 72.20370173454285

Validation Accuracy: 72.28333353996277

\*\*\*\*\* Epoch 4 \*\*\*\*\*

Training Accuracy: 70.30740976333618  
Validation Accuracy: 70.20000219345093

-----  
\*\*\*\*\* Epoch 5 \*\*\*\*\*  
Training Accuracy: 67.44258999824524  
Validation Accuracy: 67.03333258628845  
-----

\*\*\*\*\* Epoch 6 \*\*\*\*\*  
Training Accuracy: 65.14444351196289  
Validation Accuracy: 64.78333473205566  
-----

\*\*\*\*\* Epoch 7 \*\*\*\*\*  
Training Accuracy: 64.50185179710388  
Validation Accuracy: 64.16666507720947  
-----

\*\*\*\*\* Epoch 8 \*\*\*\*\*  
Training Accuracy: 64.39074277877808  
Validation Accuracy: 63.91666531562805  
-----

\*\*\*\*\* Epoch 9 \*\*\*\*\*  
Training Accuracy: 64.3240749835968  
Validation Accuracy: 63.81666660308838  
-----

\*\*\*\*\* Epoch 10 \*\*\*\*\*  
Training Accuracy: 64.25740718841553  
Validation Accuracy: 63.78333568572998  
-----

\*\*\*\*\* Epoch 11 \*\*\*\*\*  
Training Accuracy: 64.22777771949768  
Validation Accuracy: 63.733333349227905  
-----

\*\*\*\*\* Epoch 12 \*\*\*\*\*  
Training Accuracy: 64.22222256660461  
Validation Accuracy: 63.716667890548706  
-----

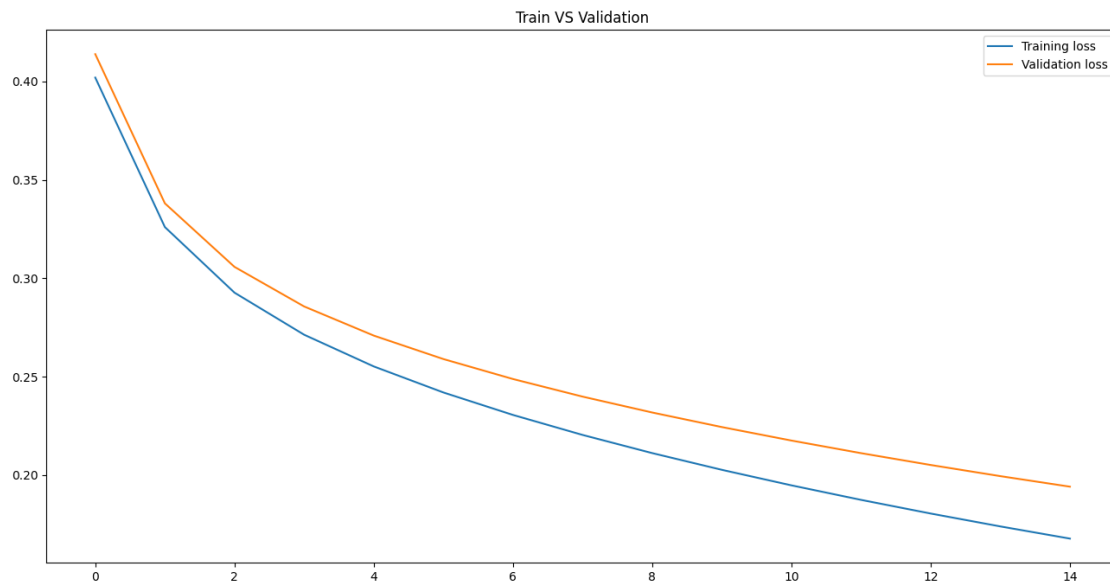
\*\*\*\*\* Epoch 13 \*\*\*\*\*  
Training Accuracy: 64.21852111816406  
Validation Accuracy: 63.733333349227905  
-----

\*\*\*\*\* Epoch 14 \*\*\*\*\*  
Training Accuracy: 64.21666741371155  
Validation Accuracy: 63.68333101272583  
-----


\*\*\*\*\* Epoch 15 \*\*\*\*\*  
Training Accuracy: 64.21296000480652  
Validation Accuracy: 63.70000243186951  
-----

Test Accuracy: 64.89999890327454

```
[ ]: plt.figure(figsize=(16,8))
plt.plot(Params[1],label='Training loss')
plt.plot(Params[2],label='Validation loss')
plt.title('Train VS Validation')
plt.legend()
plt.show()
```



```
[ ]: Mat = Confusion_Matrix(Params[-1],Y_test[0])
plot_confusion_matrix(Mat)
```



## 6 Using Pytorch

```
[ ]: class MLP(torch.nn.Module):
    def __init__(self, input_dim, output_dim):
        super().__init__()

        self.input_fc = torch.nn.Linear(input_dim, 500)
        self.hidden1_fc = torch.nn.Linear(500, 250)
        self.hidden2_fc = torch.nn.Linear(250, 100)
        self.output_fc = torch.nn.Linear(100, output_dim)

    def forward(self, x):
        batch_size = x.shape[0]
        h_1 = torch.nn.functional.relu(self.input_fc(x))
        h_2 = torch.nn.functional.relu(self.hidden1_fc(h_1))
        h_3 = torch.nn.functional.relu(self.hidden2_fc(h_2))
        y_pred = self.output_fc(h_3)
        return y_pred, h_3
```

```
[ ]: model = MLP(784,10)
```

```
[ ]: def count_parameters(model):  
    return sum(p.numel() for p in model.parameters() if p.requires_grad)
```

```
[ ]: optimizer = torch.optim.Adam(model.parameters(),lr=0.001)  
    criterion = torch.nn.CrossEntropyLoss()  
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[ ]: model = model.to(device)  
    criterion = criterion.to(device)
```

```
[ ]: def calculate_accuracy(y_pred, y):  
    top_pred = y_pred.argmax(1, keepdim=True)  
    correct = top_pred.eq(y.view_as(top_pred)).sum()  
    acc = correct.float() / y.shape[0]  
    return acc
```

```
[ ]: def train(model, iterator, optimizer, criterion, device):  
  
    epoch_loss = 0  
    epoch_acc = 0  
  
    model.train()  
  
    for data in iterator:  
        x=torch.flatten(data[0],start_dim=1)  
        y=data[1]  
        x = x.to(device)  
        y = y.to(device)  
  
        optimizer.zero_grad()  
  
        y_pred, _ = model(x)  
  
        loss = criterion(y_pred, y)  
  
        acc = calculate_accuracy(y_pred, y)  
  
        loss.backward()  
  
        optimizer.step()  
  
        epoch_loss += loss.item()  
        epoch_acc += acc.item()  
  
    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

```
[ ]: def evaluate(model, iterator, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for data in iterator:
            x=torch.flatten(data[0],start_dim=1)
            y=data[1]
            x = x.to(device)
            y = y.to(device)

            y_pred, _ = model(x)

            loss = criterion(y_pred, y)

            acc = calculate_accuracy(y_pred, y)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

```
[ ]: EPOCHS = 10

best_valid_loss = float('inf')

for epoch in range(EPOCHS):

    train_loss, train_acc = train(model, train_iterator, optimizer, criterion, ↵
    ↪device)

    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion, device)
    # To Save the Model
    if valid_loss < best_valid_loss:
        best_valid_loss = valid_loss
        torch.save(model.state_dict(), 'tut1-model.pt')

    print(f'Epoch: {epoch+1:02} ')
    print(f'\tTrain Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
    print(f'\tVal. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')
```

```

Epoch: 01
    Train Loss: 0.275 | Train Acc: 91.88%
    Val. Loss: 0.145 | Val. Acc: 95.88%
Epoch: 02
    Train Loss: 0.101 | Train Acc: 96.89%
    Val. Loss: 0.104 | Val. Acc: 96.97%
Epoch: 03
    Train Loss: 0.065 | Train Acc: 97.93%
    Val. Loss: 0.095 | Val. Acc: 97.30%
Epoch: 04
    Train Loss: 0.049 | Train Acc: 98.55%
    Val. Loss: 0.095 | Val. Acc: 97.40%
Epoch: 05
    Train Loss: 0.038 | Train Acc: 98.80%
    Val. Loss: 0.105 | Val. Acc: 97.58%
Epoch: 06
    Train Loss: 0.032 | Train Acc: 98.95%
    Val. Loss: 0.103 | Val. Acc: 97.28%
Epoch: 07
    Train Loss: 0.026 | Train Acc: 99.21%
    Val. Loss: 0.110 | Val. Acc: 97.37%
Epoch: 08
    Train Loss: 0.024 | Train Acc: 99.24%
    Val. Loss: 0.092 | Val. Acc: 97.90%
Epoch: 09
    Train Loss: 0.021 | Train Acc: 99.31%
    Val. Loss: 0.092 | Val. Acc: 97.98%
Epoch: 10
    Train Loss: 0.018 | Train Acc: 99.42%
    Val. Loss: 0.112 | Val. Acc: 97.58%

```

```
[ ]: test_loss, test_acc = evaluate(model, test_iterator, criterion, device)
```

```
[ ]: print('Test Loss: ',test_loss)
      print('Test Accuracy: ',test_acc)
```

```

Test Loss:  0.10441038757562637
Test Accuracy:  0.9795999526977539

```

```
[ ]: def get_predictions(model, iterator, device):

    model.eval()

    images = []
    labels = []
    probs = []

```



```

with torch.no_grad():

    for data in iterator:
        x=torch.flatten(data[0],start_dim=1)
        y=data[1]
        x = x.to(device)

        y_pred, _ = model(x)

        y_prob = torch.nn.functional.softmax(y_pred, dim=-1)

        images.append(x.cpu())
        labels.append(y.cpu())
        probs.append(y_prob.cpu())

    images = torch.cat(images, dim=0)
    labels = torch.cat(labels, dim=0)
    probs = torch.cat(probs, dim=0)

    return images, labels, probs

```

```
[ ]: from sklearn import metrics
```

```
[ ]: def plot_confusion_matrix(labels, pred_labels):

    fig = plt.figure(figsize=(10, 10))
    ax = fig.add_subplot(1, 1, 1)
    cm = metrics.confusion_matrix(labels, pred_labels)
    cm = metrics.ConfusionMatrixDisplay(cm, display_labels=range(10))
    cm.plot(values_format='d', cmap='Blues', ax=ax)

```


```
[ ]:
```

```
[ ]: images, labels, probs = get_predictions(model, test_iterator, device)

    pred_labels = torch.argmax(probs, 1)

```

```
[ ]: plot_confusion_matrix(labels, pred_labels)
```



```
[ ]: corrects = torch.eq(labels, pred_labels)
      incorrect_examples = []

      for image, label, prob, correct in zip(images, labels, probs, corrects):
          if not correct:
              incorrect_examples.append((image, label, prob))

      incorrect_examples.sort(reverse=True,
                             key=lambda x: torch.max(x[2], dim=0).values)
```

```
[ ]: def plot_most_incorrect(incorrect, n_images):

      rows = int(np.sqrt(n_images))
```

```

cols = int(np.sqrt(n_images))

fig = plt.figure(figsize=(20, 10))
for i in range(rows*cols):
    ax = fig.add_subplot(rows, cols, i+1)
    image, true_label, probs = incorrect[i]
    true_prob = probs[true_label]
    incorrect_prob, incorrect_label = torch.max(probs, dim=0)
    ax.imshow(image.view(28, 28).cpu().numpy(), cmap='bone')
    ax.set_title(f'true label: {true_label} ({true_prob:.3f})\n'
                f'pred label: {incorrect_label} ({incorrect_prob:.3f})')
    ax.axis('off')
fig.subplots_adjust(hspace=0.5)

```

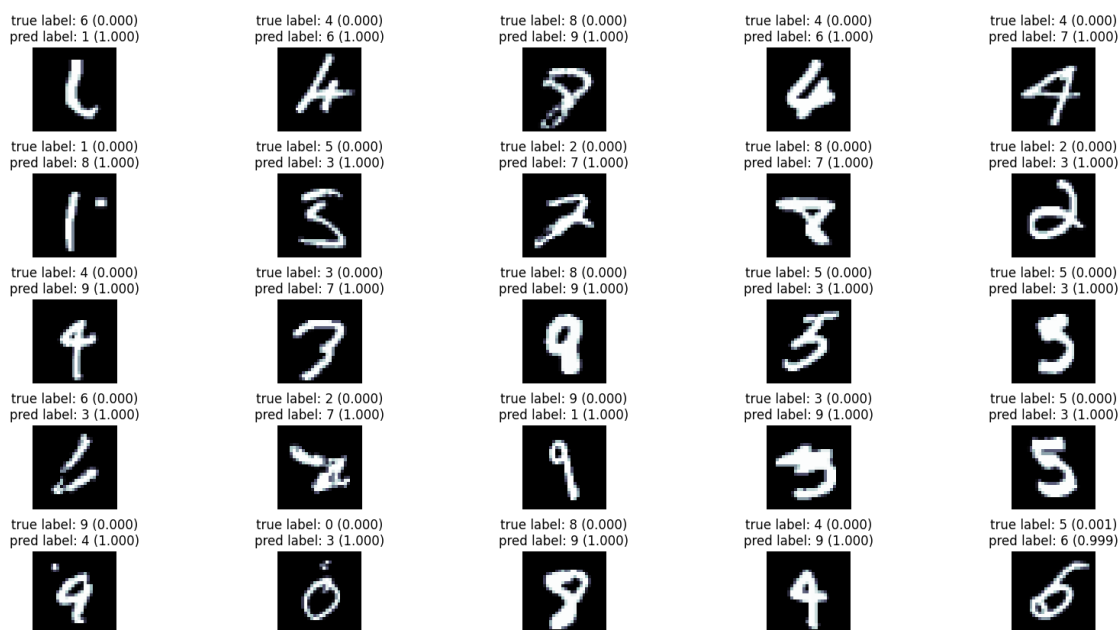
6.1 In Reality Model had done a good job as we can see the following:

```

[ ]: N_IMAGES = 25

plot_most_incorrect(incorrect_examples, N_IMAGES)

```



```

[ ]:

```

## 7 Summary

1. Sigmoid With  $Epoch = 15$  Here are the Details:  $> * Learning\ rate = 0.01$

- Train Accuracy : 67.13%
  - Validation Accuracy: 66.39%
  - Test Accuracy: 67.45%
  - Time: 4m10sec > \**Learning rate* = 0.01
  - Train Accuracy : Not Learning
  - Validation Accuracy: —
  - Test Accuracy: —
  - Time: — > \**Learning rate* = 1
  - Train Accuracy : 92.17%
  - Validation Accuracy: 91.85%
  - Test Accuracy: 92.14%
  - Time: 4m10sec
- 

2. **ReLU** With *Epoch* = 15 Here are the Details: > \* *Learning rate* = 0.01

- Train Accuracy : 90.47%
- Validation Accuracy: 90.20%
- Test Accuracy: 91.00%
- Time: 4m1sec > \**Learning rate* = 0.1
- Train Accuracy : 97.33%
- Validation Accuracy: 96.1%
- Test Accuracy: 96.64%
- Time: 3m55s > \**Learning rate* = 1
- Train Accuracy : 99.83
- Validation Accuracy: 97.75
- Test Accuracy: 98.07
- Time: 3m0sec

#Regularisation penalty=0.01 1. **ReLU** With *Epoch* = 15 Here are the Details: > \* *Learning rate* = 0.01 \* Train Accuracy : — \* Validation Accuracy: — \* Test Accuracy: — \* Time: 4m10sec 2. **Tanh** > *Learning rate* = 0.01 Train Accuracy : 65.04 \* Validation Accuracy: 63.07 \* Test Accuracy: 64.70 \* Time: — —

1. **Pytorch** With *Epoch* = 10 and Adam optimizer Here are the Details: > \* *Learning rate* = 0.01

- Train Accuracy : 97.95
  - Validation Accuracy: 95.07
  - Test Accuracy: 95.20
  - Time: 3m45s > \**Learning rate* = 0.01
  - Train Accuracy : Not doing good job
  - Validation Accuracy: —
  - Test Accuracy: —
  - Time: — > \**Learning rate* = 0.001
  - Train Accuracy : 99.48
  - Validation Accuracy: 97.85
  - Test Accuracy: 97.75
  - Time: 2m55sec
-